

Tetration

Cisco Tetration Analytics

Building better Dashboards v1

Ryan Tischer, CCIE 11459

Cisco

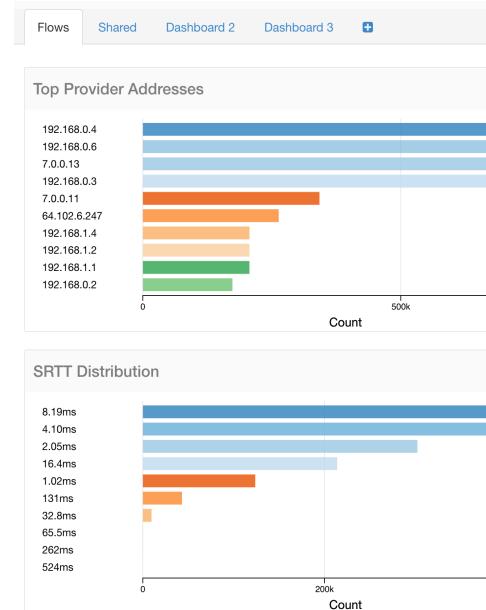
1/2017

TSA Call to Action

- Create interesting Dashboards are share.
- Better Demos that include Dashboards
- Need to figure out the best way to share? Github?
- Suggest we build a team to focus on TA Dashboards

Tetration Dashboards

- Summary of Tetration Data
- Designed for NOCs,
Application owners, or C-level
- Can be broad i.e. total BW or
based on specific data (App
specific)
- Based on Druid



What is Druid

- Column-oriented, distributed data store
- Highly scalable (Petabytes)
- Supports real-time streams (Spark/Kafka)
- **Sub-second Queries**
- **Queries built in JSON formatted data**
- Used to power TA Dashboards and Flow Search

Example Data in Druid

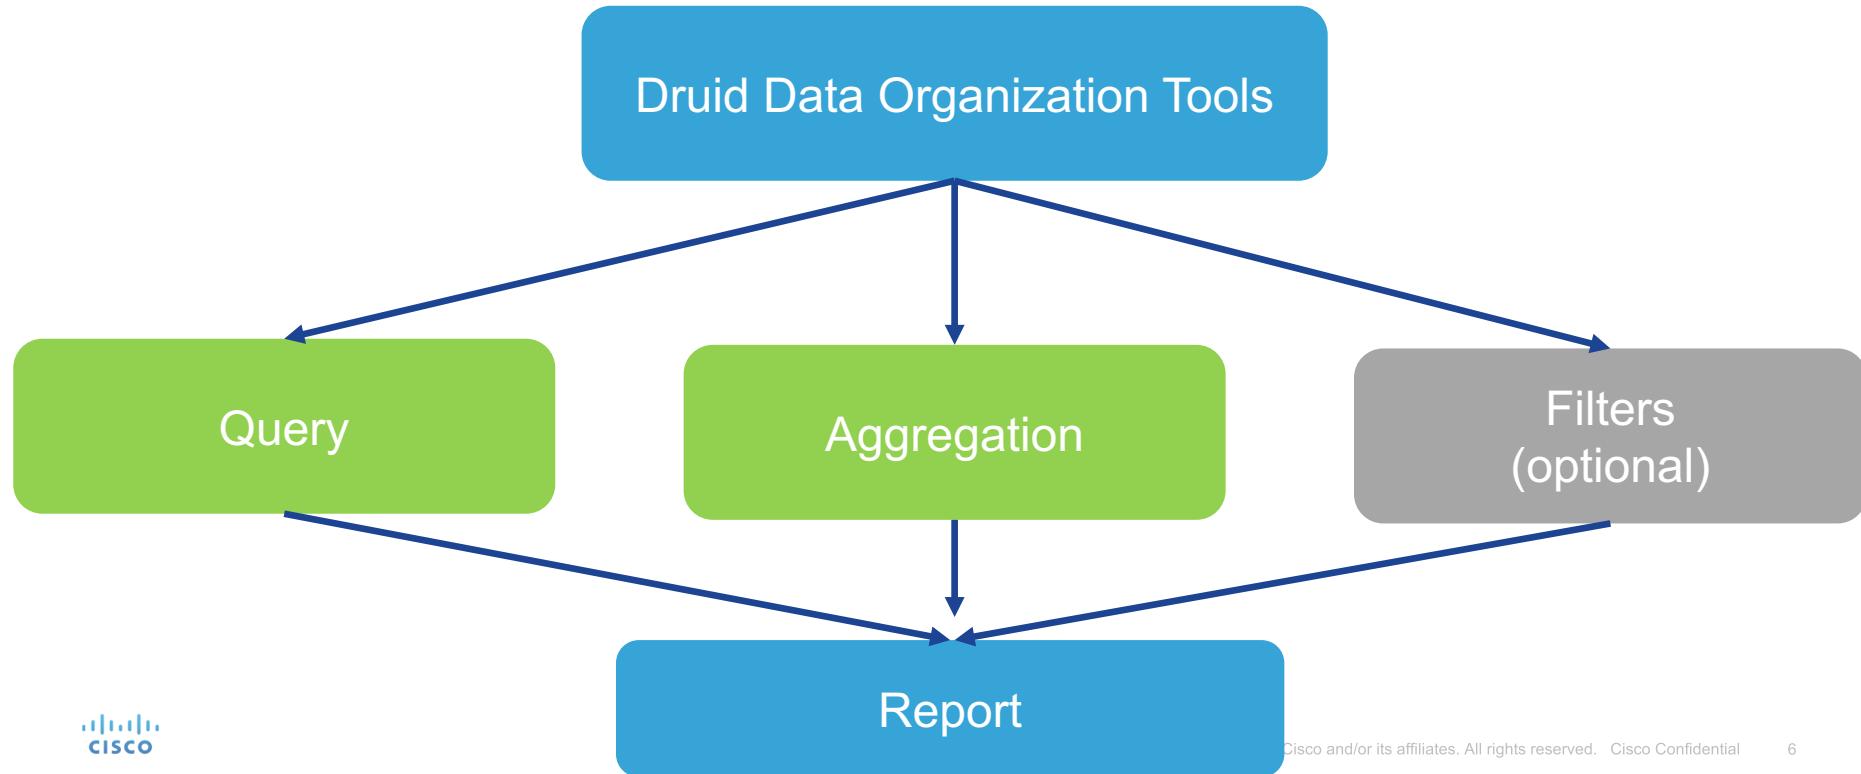
Wikipedia “edit” events

Timestamp	Page	Username	Gender	City	Characters Added	Characters Removed
2011-01-01T01:00:00Z	Justin Bieber	Boxer	Male	San Francisco	1800	25
2011-01-01T01:00:00Z	Justin Bieber	Reach	Male	Waterloo	2912	42
2011-01-01T02:00:00Z	Ke\$ha	Helz	Male	Calgary	1953	17
2011-01-01T02:00:00Z	Ke\$ha	Xeno	Male	Taiyuan	3194	170

Time stamped events in a table

Column-oriented

The Magic of Druid (TA Dashboards) is organizing data in away that answers questions.



Querying Data in Druid

- Three types of Queries
 - Timeseries
 - TopN
 - GroupBy
- First of two required inputs (Aggregation also required)
- Other query types are permitted, discussed later

Time series Query

- Group in blobs of time, formatted in JSON
- User set granularity (minute, hour, year, 15 minutes, etc)
- Druid returns a data point for each interval.
- For example a timeseries query by hour would produce

```
[ { "Count": 1511410, ← The thing were counting  
  "timestamp": "2017-01-16T17:00:00.000Z" }, ← The timestamp (by hour)  
  { "Count": 111644,  
    "timestamp": "2017-01-16T18:00:00.000Z" } ]  
  Repeat for the next hour
```

The source of the Data is discussed later

TopN Query

- Sorted results of a data source, formatted in JSON
- Example Top talkers
- Druid returns a data point for each dimension (talker)
- For example top talkers by IP address

```
[ { "count": 73137, ← The thing were counting  
  "src_address": "1.1.11.15", ← Identifier  
  "timestamp": "2017-01-16T17:03:23.000Z" },  
  { "count": 72888, ← Repeat for next top talker  
  "src_address": "1.1.11.10",  
  "timestamp": "2017-01-16T17:03:23.000Z" },
```

Groupby Query

- Multi-dimension query, formatted in JSON
- Similar to SQL “Group by”
- Druid returns a data point for each object
- For example source address + vrf + src port

```
{ "src_port": "1065",  
  "vrf_id": "1",  
  "Received bytes": 3024,  
  "src_address": "7.0.0.32",  
  "timestamp": "2017-01-16T17:00:00.000Z" },
```

Grouped Object Identifier

The thing were counting

Repeat for next object

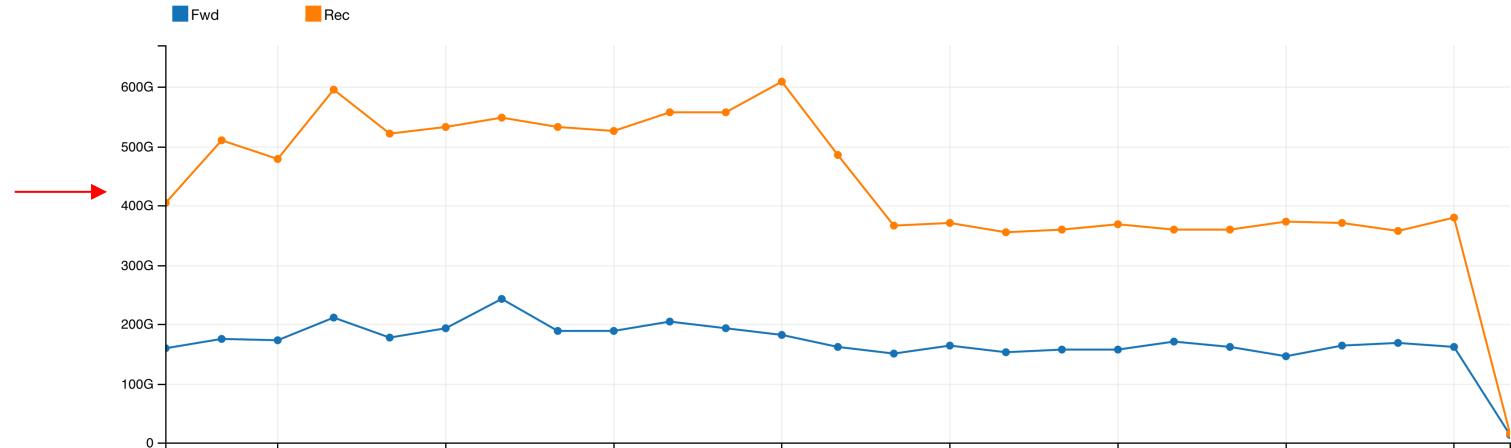
Aggregations

- Aggregate or group data into a single row
- Used with Queries
- For example
 - Query data in X timeseries, by hour and aggregate by Y types/metric
 - Query data from 1pm - 2pm by minutes and aggregate by sent bytes
- Metric includes count or any valid data within the dataset
- Multiple Aggregations Supported
 - Must be like data (sent/receive **bytes**) for the x axis to make sense

Query + Aggregations

- Queries build the X axis (per data source)
- Aggregations build the Y axis

Aggregation



Aggregations types / metrics

- Types
 - Count - Count the number of rows
 - DoubleSum – 64 bit integer value. Requires a metric
 - LongSum – 64 bit floating point value. Requires a metric. Not commonly used
 - Min/Max – Minimum or Maximum value in the set
 - Other are available, discussed later

Druid Filters

- Optional filter which rows are included in the dataset.
- Equivalent for SQL WHERE clause
- For example find rows where host name = SQL
- Two basic types of filters. Both are **CASE SENSITIVE**
 - Selector
 - Match a specific dimension. Hostname = App01
 - Regex
 - Any Regular expression (JAVA)
 - **Regex is preferred.**
- Other options discussed later

Regex Cheat Sheet

- “*” match zero or more characters
 - “App*” matches: App, ”App1”, “App2” , “App01”, “Appasfadsfasdfasab”
- “+” match one or more characters
 - App+ matches: , ”App1”, “App2” , “App01”, “Appasfadsfasdfasab”
- (?i) before an expression makes the search **case-insensitive**
 - (?i)app01 matches: “app01”, “apP01”, “ApP01”, “App01”, etc
- [...] matches a range
 - App0[1-10] matches: “App01”, “App09” etc
 - Ap[a-Z]01 matches: “Apt01”, “App01”, “ApP01”, etc
- “^” is do not match
 - App0[^1] will not match App01

“\$” signals end of string
App01\$ matches: App01
(best practice)

Building Dashboards in Tetration

Time Frame

Jan 16 1:25pm - Jan 16 2:25pm ▾

Back to Dashboard

New Chart

Chart Name

Chart Name

Data Source

Aggregations
Aggregations

Type

-- select type --

+ Aggregation

Filters

Optional Filters
Only 1 supported

JSON Request

```
{"intervals":["2017-01-16T13:25:06-06:00/2017-01-16T14:25:06-06:00"]}
```

This holds the json formatted query. May be edited by hand

Refresh Chart

Save Chart

Download Data as CSV

SAVE

TA Data Sources

- Flow
- Hosts
- Attacks (no data yet)
- Policy Delta (no data yet)
- ADM runs appear as choices
 - However have no data (maybe future)
- Would be super nice if we had a dump of the raw table

TA Data Source - Flow

- All flow data collected
- Exposed in flow search
- Metrics on any number
- Filter on anything
- Other data may be available

```
1 [
2   "src_hostname": "web01",
3   "src_address": "7.0.0.41",
4   "src_port": "49167",
5   "dst_hostname": "tet-13.internal",
6   "dst_address": "172.26.46.35",
7   "dst_port": "5640",
8   "proto": "TCP",
9   "key_type": "1",
10  "start_timestamp": "1484512020000000",
11  "rev_process_string": "/opt/tetration/collector/tet-collector --config_file /etc/tetration
12    /collector/collector.config --timestamp_flow_info --logtostderr --max_num_ssl_sw_sensors 63000
13    --enable_client_certificate true --write_empty_files true",
14  "rev_process_owner": "tetter",
15  "byte_count_dim": "16384",
16  "pkt_count_dim": "32",
17  "flow_duration_dim": "68719476736",
18  "data_duration_dim": "68719476736",
19  "srtt_dim_usec": "16384",
20  "network_latency_dim_usec": "0",
21  "server_app_latency_dim_usec": "0",
22  "vrf_name": "Default",
23  "fwd_policies": "MAYBE_PERMITTED:UNKNOWN",
24  "rev_policies": "MAYBE_PERMITTED:UNKNOWN",
25  "fwd_tags": "PSH ACK CWR",
26  "rev_tags": "ACK",
27  "record_start_ts_usec": 1484580180000000,
28  "record_end_ts_usec": 1484580239372603,
29  "vrf_id": "1",
30  "count": 1,
31  "srtt_usec": 16875,
32  "total_network_latency_usec": 0,
33  "server_app_latency_usec": 0,
34  "fwd_pkts": 41,
35  "rev_pkts": 41,
36  "fwd_bytes": 19984,
37  "rev_bytes": 2238,
38  "fwd_flow_duration": 68219372603,
39  "rev_flow_duration": 68219372593,
40  "fwd_data_duration": 68219372603,
41  "rev_data_duration": 0,
42  "timestamp": "2017-01-16T15:23:00.000Z"
43 ]
```

TA Data Source - Hosts

- All host data collected
- Other data may be available
- Aggregations →
- Filter on Host data
- All attacker/victim data is 0. Future release?



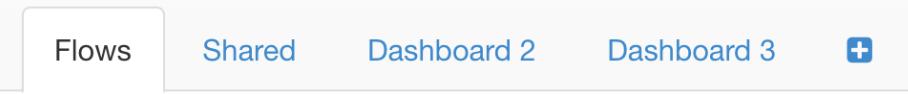
attacker_count
count
ddos_intensity
reputation
sum_bytes_received
sum_bytes_sent
sum_packets_received
sum_packets_sent
sum_received_fin
sum_received_syn
sum_sent_fin
sum_sent_syn
victim_count
volumetric_ddos_debug

GUI → JSON

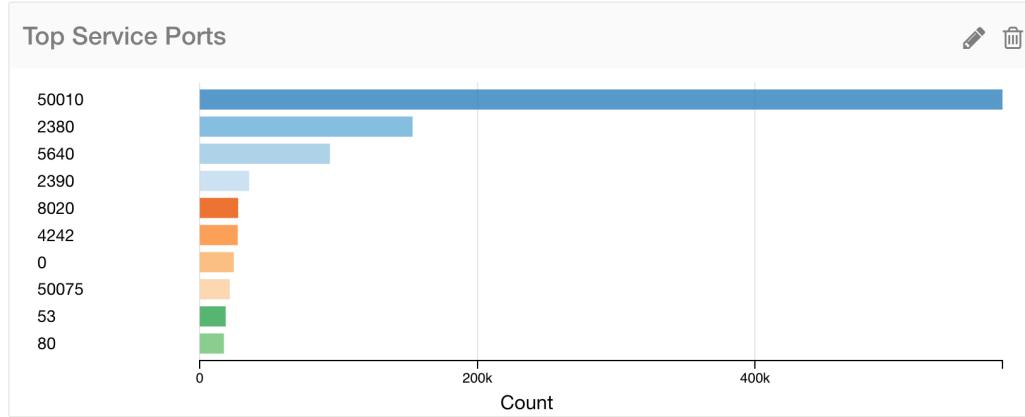
- The GUI builds the JSON request in real-time
- Can be edited by hand (more on this later)
- Used to share queries between dashboards or systems
- Use <http://www.jsoneditoronline.org/> or other JSON editor to help visualize data
- Example →

```
{"dataSource":"flows",
  "granularity":"all",
  "queryType":"topN","threshold":10,"metric":"count",
  "aggregations":[{"type":"count","fieldName":"","name":"count"}],
  "intervals":["2017-01-16T13:25:06-06:00/2017-01-
  16T14:25:06-06:00"],"dimension":"dst_port"}
```

Dashboard hints

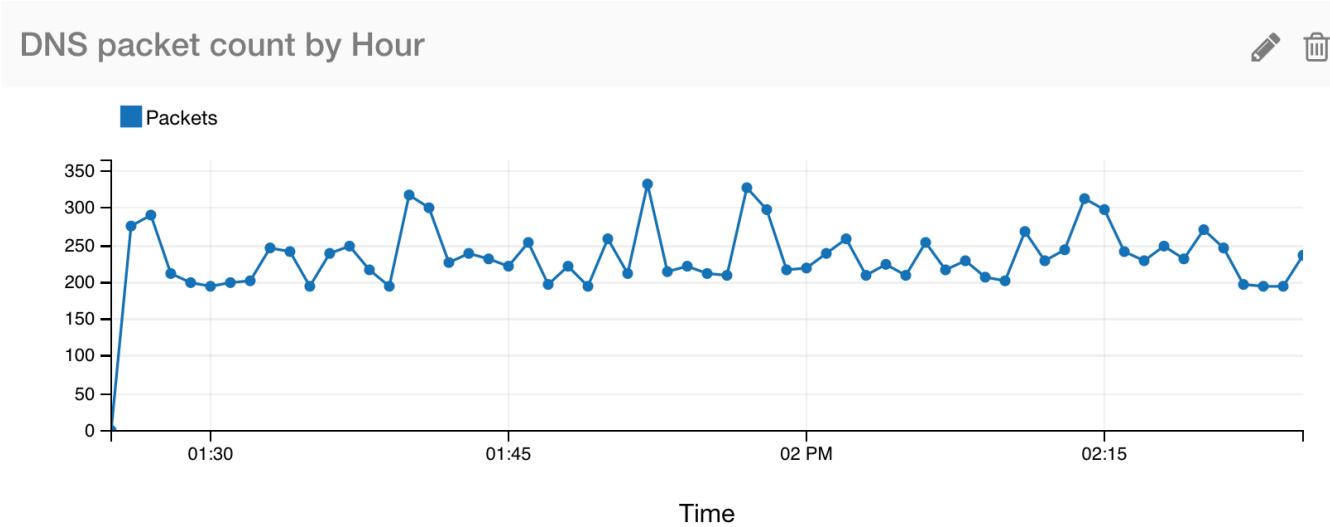
- Chart options will change based on...
 - Data Source
 - Query Type
 - Aggregations
- Use Chart preview and **JSON Response** to test data
 - If data count is all 0s you did something wrong
- TA Dashboard Tabs 
 - Flows tab is built in and can not be changed
 - Shared tab is shared between all users. Can be changed. Use as example
 - Users create personal dashboards

Dashboard Examples – Top Service Ports



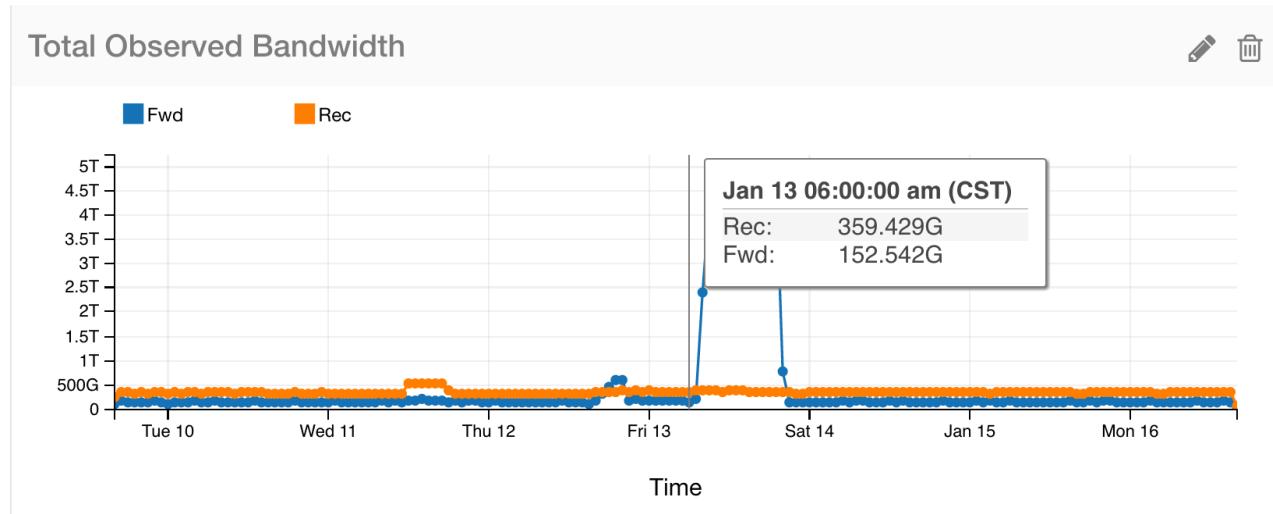
```
{"dataSource":"flows","granularity":"all","queryType":"topN","threshold":10,"metric":"count","aggregations":[{"type":"count","fieldName":"","name":"count"}],"intervals":["2017-01-16T13:25:06-06:00/2017-01-16T14:25:06-06:00"],"dimension":"dst_port"}
```

Dashboard Examples – DNS Packet Count



```
{"intervals":["2017-01-16T13:25:06-06:00/2017-01-16T14:25:06-06:00"], "dataSource": "flows", "granularity": "minute", "queryType": "timeseries", "aggregations": [{"type": "doubleSum", "fieldName": "rev_pkts", "name": "Packets"}], "filter": {"type": "selector", "dimension": "dst_port", "value": "53"}}
```

Multiple Aggregations



Aggregations

Type	Name	Metric
doubleSum	Fwd	fwd_bytes
Type	Name	Metric
doubleSum	Rec	rev_bytes

TA Dashboard Off-Roading

Dashboard Off-roading

- Druid provides a rich query/visualization tools that are not exposed by the GUI
- Require editing JSON by hand
- A lot more work required!

Filters with AND / OR Expressions

- Druid Filters support AND/OR/NOT expressions
- Example...

Expression

Regular Filter

```
    "filter": {  
        "type": "and",  
        "fields": [  
            {  
                "type": "regex",  
                "dimension": "dst_hostname",  
                "pattern": "[Ss][qQ][Ll]*"  
            },  
            {  
                "type": "selector",  
                "dimension": "dst_port",  
                "value": "1433"  
            }  
        ],  
        "dimension": null  
    }
```

Filters with NOT Expression

- Example...

```
1  {  
2      "dataSource": "flows",  
3      "granularity": "all",  
4      "queryType": "topN",  
5      "threshold": 10,  
6      "metric": "count",  
7      "aggregations": [  
8          {  
9              "type": "count",  
10             "fieldName": "",  
11             "name": "count"  
12         }  
13     ],  
14     "intervals": [  
15         "2017-01-31T13:36:45-06:00/2017-01-31T14:36:45-06:00"  
16     ],  
17     "dimension": "proto",  
18     "filter": {  
19         "type": "not",  
20         "field": {  
21             {  
22                 "type": "regex",  
23                 "dimension": "proto",  
24                 "pattern": "TCP"  
25             }  
26         }  
27     }  
}
```

JavaScript Functions

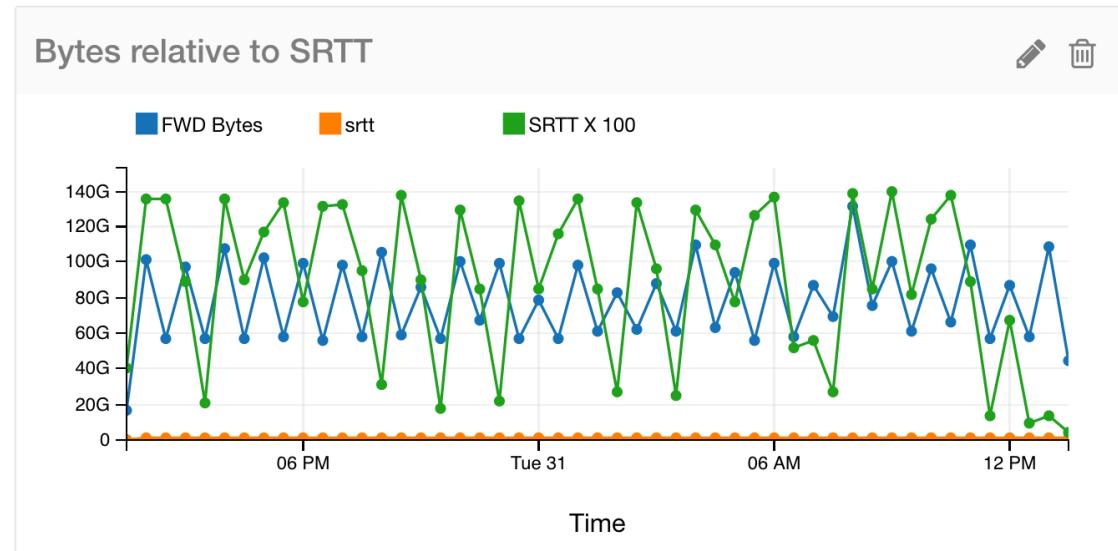
- Javascript functions serve as aggregators, filters, or post-aggregations. **Must return a number**
- Useful for complex manipulation
- Druid documentation is terrible and no one seems to use this
- JS functions will provide much better data and add context
- Examples include
 - Ratio of traffic behind load balancers
 - Total resources consumed for a multitier application
 - Port BW consumption / total
 - TCP sessions / server capacity ← app scaling

```

1 [
2   "intervals": [
3     "2017-01-30T13:50:05-06:00/2017-01-31T13:50:05-06:00"
4   ],
5   "dataSource": "flows",
6   "granularity": "thirty_minute",
7   "queryType": "timeseries",
8   "aggregations": [
9     {
10       "type": "doubleSum",
11       "name": "FWD Bytes",
12       "fieldName": "fwd_bytes"
13     },
14     {
15       "type": "doubleSum",
16       "name": "srtt",
17       "fieldName": "srtt_usec"
18     },
19     {
20       "type": "javascript",
21       "name": "SRTT X 100",
22       "fieldNames": [
23         "srtt_usec"
24       ],
25       "fnAggregate": "function(current, a) { return current + (a * 100); }",
26       "fnCombine": "function(partialA) { return partialA; }",
27       "fnReset": "function() { return 0; }"
28     }
29   ]
30 ]

```

JavaScript Aggregator Example



```
1 [
2   "intervals": [
3     "2017-01-30T13:50:05-06:00/2017-01-31T13:50:05-06:00"
4   ],
5   "dataSource": "flows",
6   "granularity": "thirty_minute",
7   "queryType": "timeseries",
8   "aggregations": [
9     {
10       "type": "doubleSum",
11       "name": "FWD Bytes",
12       "fieldName": "fwd_bytes"
13     },
14     {
15       "type": "doubleSum",
16       "name": "srtt",
17       "fieldName": "srtt_usec"
18     },
19     {
20       "type": "javascript",
21       "name": "SRTT X 100",
22       "fieldNames": [
23         "srtt_usec"
24       ],
25       "fnAggregate": "function(current, a) { return current + (a * 100); }",
26       "fnCombine": "function(partialA) { return partialA; }",
27       "fnReset": "function() { return 0; }"
28     }
29   ]
30 ]
```

JavaScript Aggregator Example

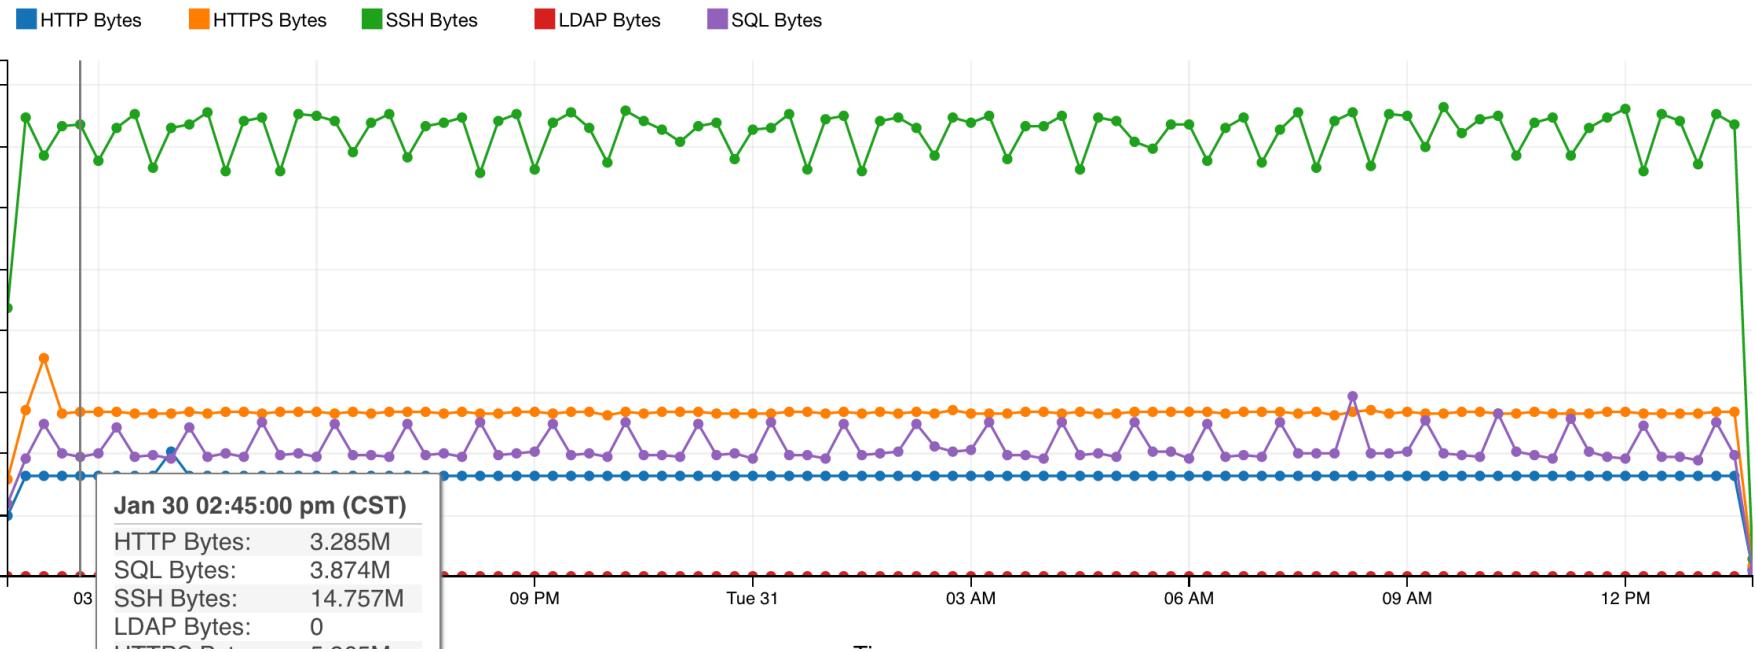
All 3 functions must be present
fnAggregate

- “current” – unclear why
- “a” – ordered list of from “fieldnames” list
- **return the result**

fnCombine - aggregate the different rows
fnReset – init value

Filtered Aggregator

- Filter aggregator on the fly



Filtered Aggregator

```
[{"intervals": [ "2017-01-30T13:50:05-06:00/2017-01-31T13:50:05-06:00" ], "dataSource": "flows", "queryType": "timeseries", "granularity": "fifteen_minute", "aggregations": [ { "type": "filtered", "name": "HTTP Bytes", "filter": { "type": "selector", "dimension": "dst_port", "value": "80" }, "aggregator": { "type": "doubleSum", "fieldName": "rev_bytes", "name": "HTTP Bytes" } }, { "type": "filtered", "name": "HTTPS Bytes", "filter": { "type": "selector", "dimension": "dst_port", "value": "443" }, "aggregator": { "type": "doubleSum", "fieldName": "rev_bytes", "name": "HTTPS Bytes" } }, { "type": "filtered", "name": "SSH Bytes", "filter": { "type": "selector", "dimension": "dst_port", "value": "22" } }
```

Post Aggregations

- TA seems to ignore input

This requires
a lot more
work and is
the focus on
v2 of this doc

Cardinality Aggregations

- Determine the number of distinct X
- This may be useful for
 - Number of servers with port 80 open
 - Number of different servers talking to SQL
 - Distinct EPGs

This requires
a lot more
work and is
the focus on
v2 of this doc